

Diploma work

Sergei Chulanov

Task description

Variant 3

Using API

<https://covidtracker.bsg.ox.ac.uk/about-api> get all data about "By country over time" for current year.

Store it into your DB: date_value, country_code, confirmed, deaths, stringency_actual, stringency.

Output the data by country_code (the country_code is set) in form of a table and sort them by deaths (or date_value which is the same) in ascending order.

| Criteria | Related Module |
|------------------------------|----------------------------------|
| SCM (Source code management) | git |
| Language | golang (front and api) |
| Tests | golang tests |
| Quality gate | golangci-lint, tfLint, sonarqube |
| Vulnerability | trivy |
| IaC | terraform |
| Orchestration | eks |
| Logging | cloudwatch |
| Monitoring | cloudwatch |
| Runtime/Deployment | gitlab-ci (test and prod) |
| Scalability/redundancy | eks (hpa) |
| Cloud and Cost efficiency | infracost |

Application description

Name: goCovid

Language: golang (go version go1.17.6 linux/amd64)

Database: MariaDB (Amazon RDS)

3-tire application (front-end, back-end, database)

back-end:

- retrieve a portion of data from API (see in your Variant) and store it in a database
- update data on demand
- update DB schema if needed on app's update

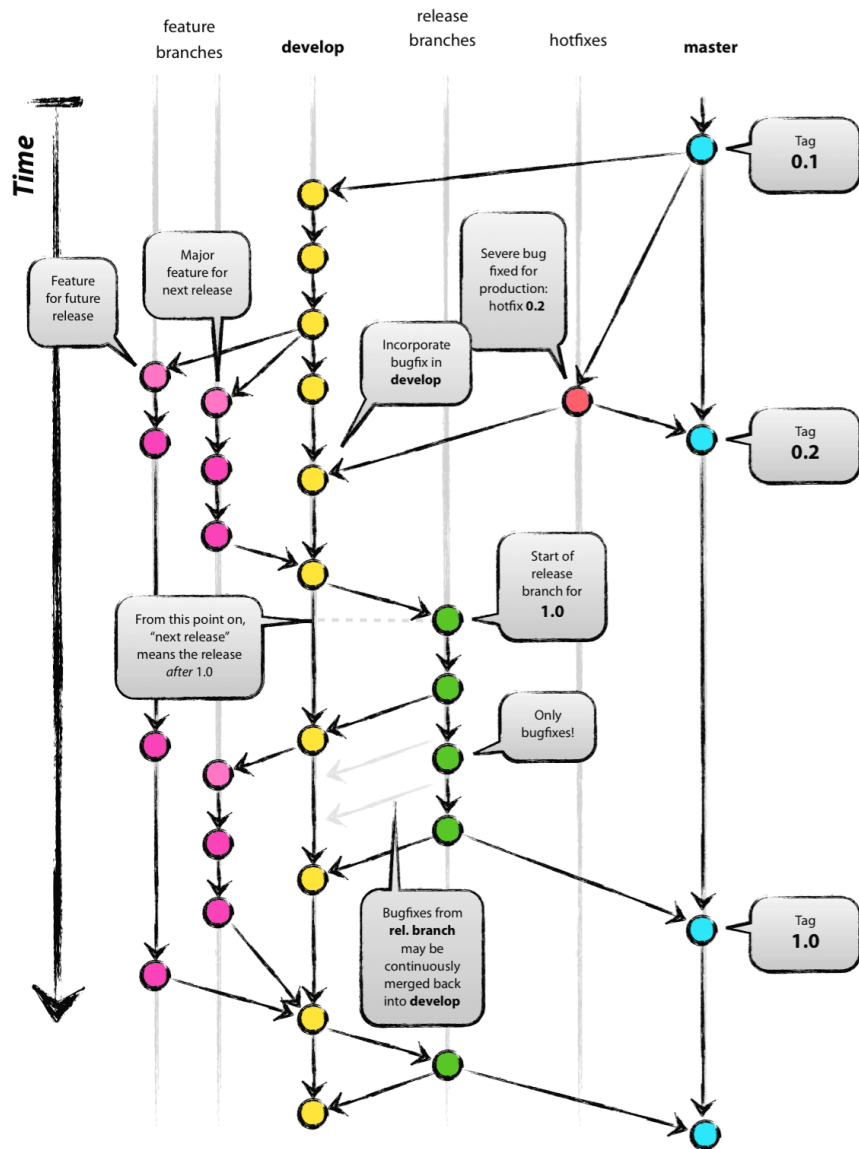
front-end:

- display any portion of the data stored in the DB
- provide a method to trigger data update process

Tools

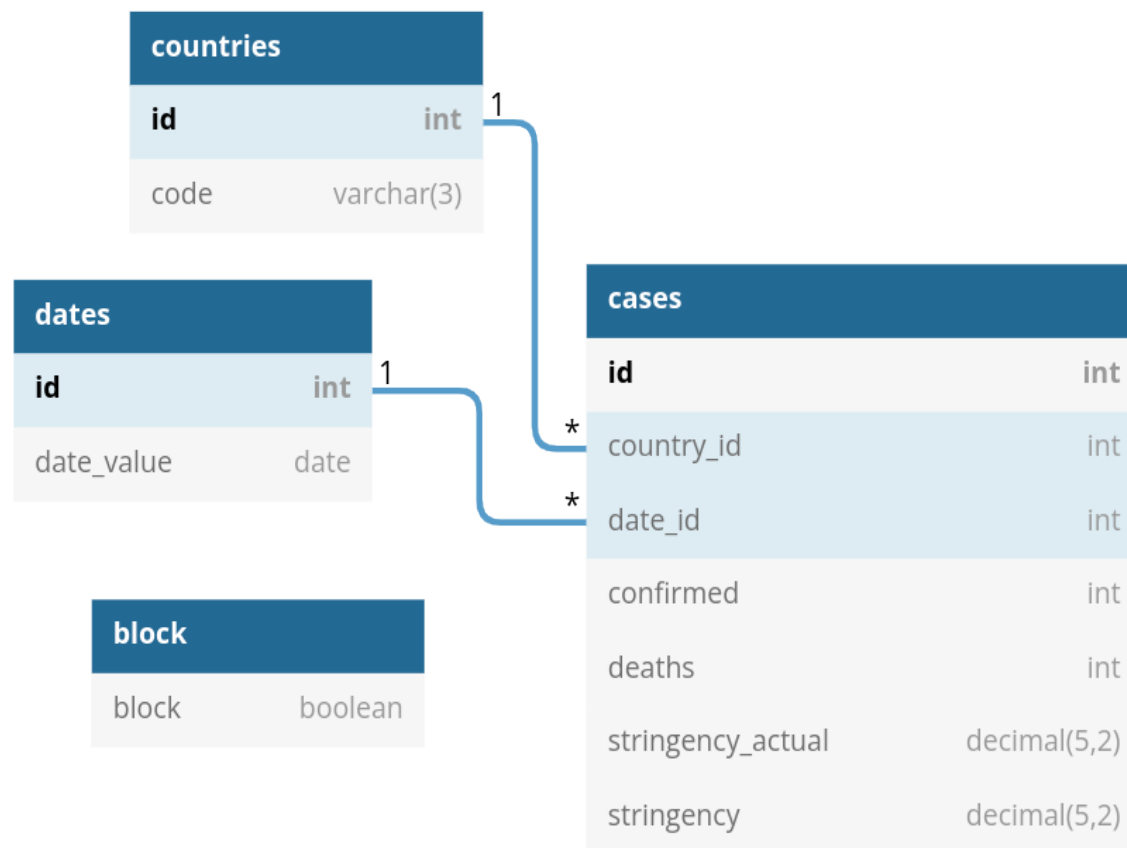
- Git
- Make
- Docker
- Terraform
- Ansible
- Helm
- EKS
- Gitlab-ci
- Infracost
- Cloudflare
- AWS
- dbdiagram.io
- app.diagrams.net

Source code management. Git. Branching strategy.



- <https://nvie.com/posts/a-successful-git-branching-model/>
- **Master**
- **Develop**
- Feature branches
 - From develop
 - Merge to develop
- Release branches
 - From develop
 - Merge into master and develop
 - Name: release-*
- Hotfix branches
 - From master
 - Merge into master and develop
 - Name: hotfix-*

Database schema



- <https://dbdiagram.io/d/61a34baa8c901501c0d5ba2c>

Table countries as C {
id int [pk, increment] // auto increment
code varchar(3)
}

Table cases {
id int [pk, increment] // auto increment
country_id int
date_id int
confirmed int
deaths int
stringency_actual decimal(5,2)
stringency decimal(5,2)
}

Table dates as DS {
id int [pk, increment] // auto increment
date_value date
}

Table block {
block boolean
}

Ref: cases.country_id > C.id

Ref: cases.date_id > DS.id

Backend (API)

```
/app # ls /app/configs/config.yml
```

```
/app/configs/config.yml
```

```
/app # ./api -help
```

Usage of ./api:

-addr string

API HTTP address (default "0.0.0.0:7000")

-fill

If it is true it fills db

```
/app # curl "http://localhost:7000/v1/help"
```

Hostname: 4abd5ea319b6

Examples:

```
curl -D -s -X GET "http://localhost:7000/v1/health-check"
```

```
curl -D -s -X GET "http://localhost:7000/v1/help"
```

```
curl -D -s -X GET "http://localhost:7000/v1/state" (Block)
```

```
curl -D -s -X GET "http://localhost:7000/v1/update" (Readiness)
```

```
curl -D -s -X GET "http://localhost:7000/v1/refresh_data"
```

```
curl -D -s -X GET "http://localhost:7000/v1/data?"
```

```
countryCode=RUS&&dateFrom=2022-01-01&&dateTo=2022-01-09&&sortBy=deaths"
```

```
SELECT
```

```
countries.code
```

```
,dates.date_value
```

```
,cases.deaths
```

```
,cases.confirmed
```

```
FROM
```

```
cases INNER JOIN countries
```

```
ON countries.id = cases.country_id INNER JOIN dates
```

```
ON dates.id = cases.date_id
```

```
WHERE
```

```
countries.code = 'USA'
```

```
AND date_value BETWEEN '2022-01-01' AND '2022-02-02'
```

```
ORDER BY
```

```
cases.deaths ASC
```



Update data



Apply

Frontend

- Generate html by go lang
- Go routine to run task on backend
- JavaScript to run code on the client side
- SessionStorage
- Twitter Bootstrap

/app # ./web -help

Usage of ./web:

-addr string

HTTP address (default "localhost:4000")

-apiVers string

API version (default "v1")

-hostname string

API hostname (default "localhost")

-port string

API port (default "8080")

Docker build with Makefile

```
$ cat Makefile.Docker
```

```
.DEFAULT_GOAL := all
```

```
# the same that .DEFAULT_GOAL for older version of  
make (<= 3.80)
```

```
# .PHONY: default
```

```
# default: clean;
```

```
.PHONY: all
```

```
all: build run
```

```
.PHONY: build
```

```
build:
```

```
$(info -> build:)
```

```
go mod tidy
```

```
mkdir -p ./bin
```

```
GOARCH=${GOARCH} GOOS=${GOOS} go build -o ./bin/$  
{BINARY_NAME}-linux ${PATH_TO_SOURCE}
```

```
.PHONY: run
```

```
run:
```

```
$(info -> run:)
```

```
./bin/${BINARY_NAME}-linux
```

```
.PHONY: clean
```

```
clean:
```

```
$(info -> clean:)
```

```
go clean
```

```
rm ./bin/${BINARY_NAME}-linux
```

```
.PHONY: messages
```

```
messages:
```

```
$(info test info message)
```

```
$(warning test warning message)
```

```
$(error test error message)
```

Docker multi-stage build with Makefile

```
FROM golang:1.17.6-alpine3.15 AS build
```

```
WORKDIR /app
```

```
ARG CGO_ENABLED=0 \  
GOARCH=amd64 \  
GOOS=linux \  
BINARY_NAME=api \  
PATH_TO_SOURCE=./cmd/api
```

```
RUN apk --no-cache add make
```

```
COPY go.mod ./\  
COPY go.sum ./\  
RUN go mod download
```

```
COPY ./app/\  
RUN make -f Makefile.Docker build
```

```
FROM alpine:3.15
```

```
ENV DB_HOSTNAME=localhost\  
ENV MIGRATIONS_DIR=/app/migrations
```

```
RUN apk --no-cache add curl
```

```
COPY --from=build /app/bin/api-linux /app/api\  
COPY --from=build /app/migrations /app/migrations\  
COPY --from=build /app/configs /app/configs
```

```
WORKDIR /app
```

```
EXPOSE 5000
```

```
CMD ["/app/api"]
```

Docker compose

```
version: '3.1'
services:

  db:
    image: mariadb:10.7.1
    environment:
      - MARIADB_DATABASE=$MARIADB_DATABASE
      - MARIADB_USER=$MARIADB_USER
      - MARIADB_PASSWORD=$MARIADB_PASSWORD
      - MARIADB_RANDOM_ROOT_PASSWORD=yes
    ports:
      - 3306:3306
    healthcheck:
      test: ["CMD", "mysql", "$MARIADB_DATABASE", "-h", "localhost", "-P", "3306", "-u",
"$MARIADB_USER", "-p$MARIADB_PASSWORD", "-e", "SELECT 1"]
      interval: 3s
      timeout: 3s
      retries: 3
      start_period: 1s

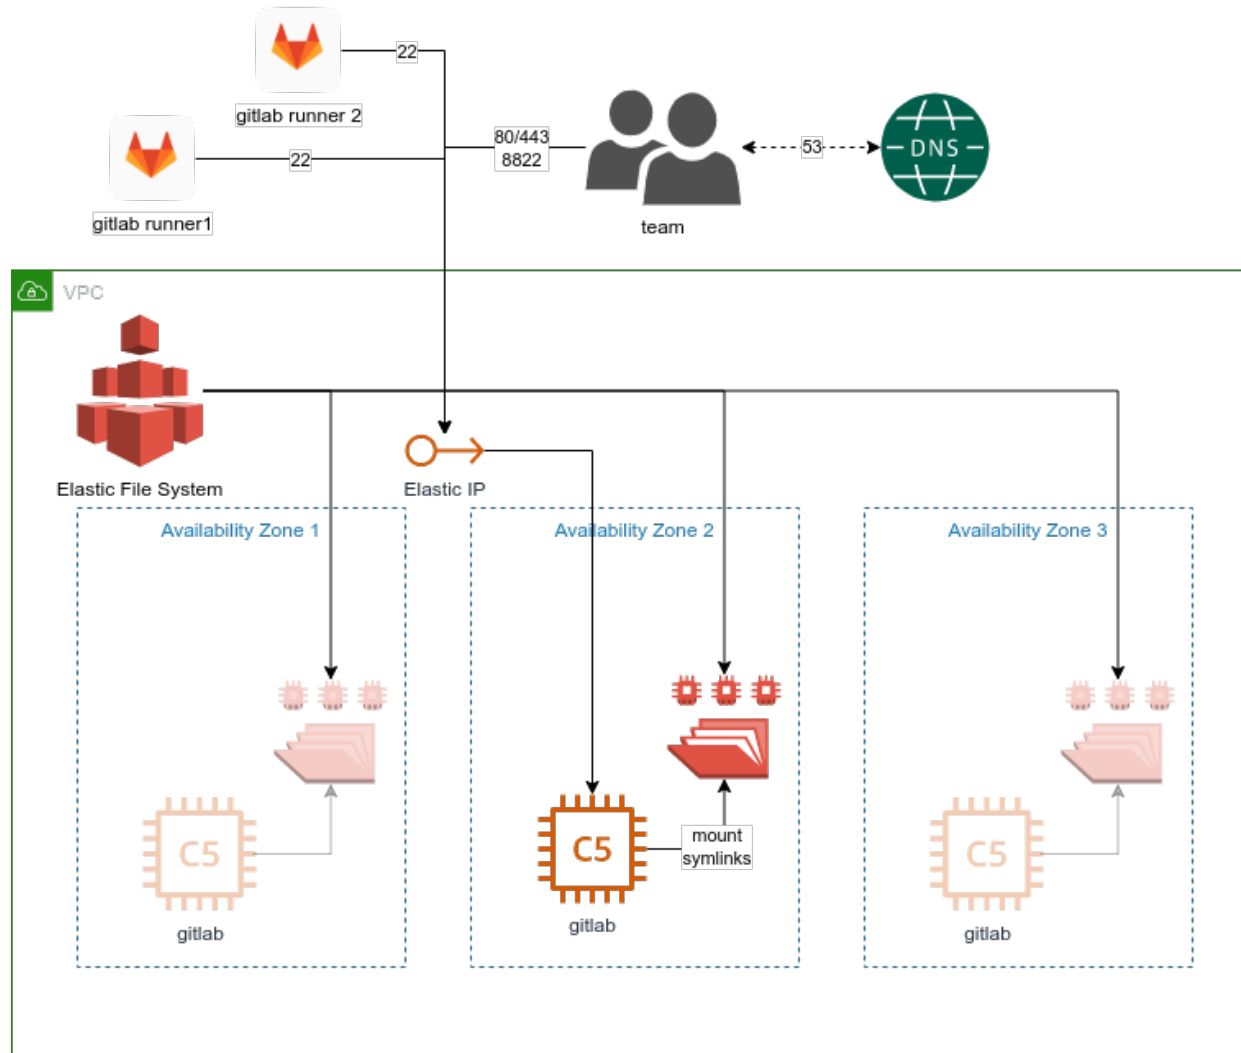
  goose:
    depends_on:
      db:
        condition: service_healthy
    build: app/backend/migrations
    links:
      - db:mysql
    environment:
      - GOOSE_DRIVER=mysql
    GOOSE_DBSTRING=$MARIADB_USER:$MARIADB_PASSWORD@tcp(mysql:3306)/$MARIADB_DATABASE
    command: ["goose", "up"]

  api_data_init:
    depends_on:
      - goose
    build: app/backend
    links:
      - db:mysql
    command: ["/app/api", "-fill=true"]

  api:
    depends_on:
      db:
        condition: service_healthy
    build: app/backend
    ports:
      - 8080:7000
    links:
      - db:mysql
    environment:
      - APP_ENDPOINT=0.0.0.0:7000
    command: ["/app/api"]

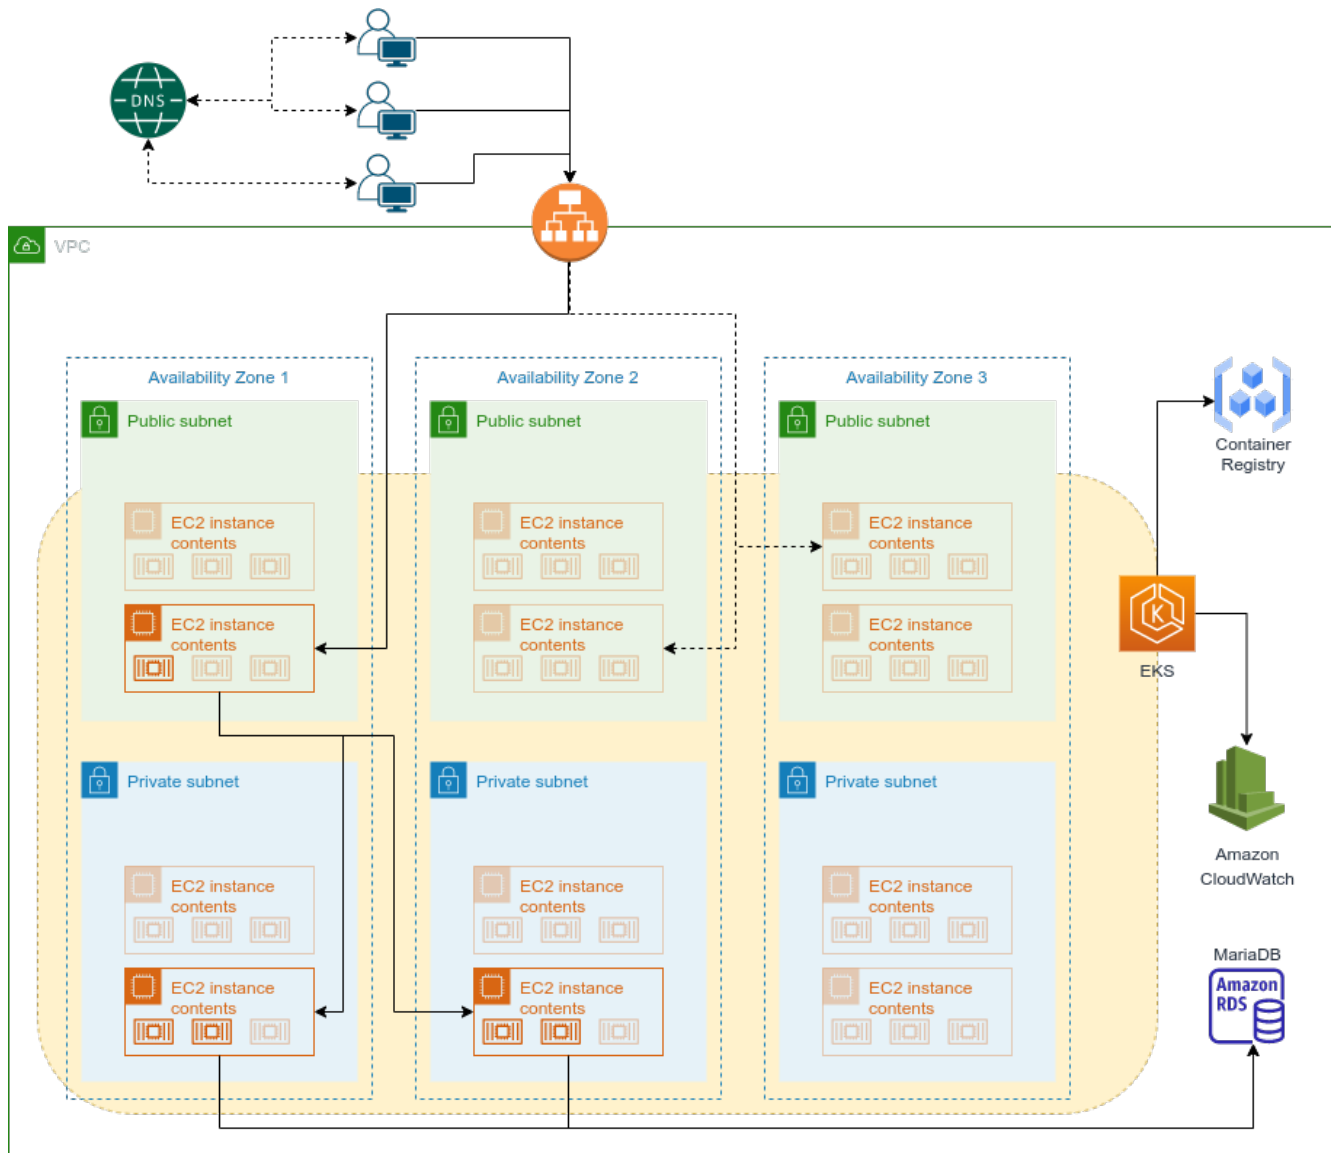
  web:
    build: app/frontend
    ports:
      - 4000:4000
    links:
      - api:api
    command: ["/app/web", "-addr", "0.0.0.0:4000", "-port", "7000", "-hostname", "api"]
```

Omnibus GitLab



- Terraform
 - cloudflare_record
 - template_file
 - EFS
 - Security groups
 - VPC
 - aws_instance
- Ansible
 - Nginx
 - Letsencrypt
 - Docker
 - EFS mount
 - -e "remove_gitlab_data=true"
- gitlab-ci
- Terraform state on gitlab
- Env variables (test and prod)
- Helm packages

Elastic Kubernetes Service



Terraform

- K8s
 - Cloudflare (with alb)
 - EKS Cluster
 - Policies
 - Amazon RDS (MariaDB)
 - Elastic container registry
 - Modules (openid_connect_provider)
 - etc
- Extra k8s
 - Autoscaler
 - Fluent-bit
 - Metrics-server
 - Container-insights
 - Alb controller
 - etc

Helm and manifests

```
$ ls helm/gocovid/  
Chart.yaml conf templates values.yaml  
$ ls helm/gocovid/conf  
config.yaml (1)  
$ ls helm/gocovid/templates  
api_deploy.yaml api_hpa.yaml api_secret_data_config.yaml api_service.yaml  
front_deploy.yaml front_hpa.yaml front_service.yaml goose_secrets.yaml  
_helpers.tpl ingress.yaml job_fill.yaml namespace.yaml  
$ ls helm/gocovid/values.yaml  
helm/gocovid/values.yaml  
$ cat gocovid/templates/api_secret_data_config.yaml  
apiVersion: v1  
kind: Secret  
metadata:  
  creationTimestamp: null  
  name: {{ .Values.api.secretName }}  
  namespace: {{ include "namespace" . }}  
type: Opaque  
data:  
  config.yaml: {{ include "apiConfig" . }} (4)
```

```
$ cat helm/gocovid/conf/config.yaml (2)  
migration_dir: {{ .Values.api.migrationDir }}  
app_endpoint: {{ .Values.api.endpoint }}  
db_host: {{ .Values.dbSettings.endpoint }}  
db_name: {{ .Values.dbSettings.name }}  
db_port: {{ .Values.dbSettings.port }}  
db_user: {{ .Values.dbSettings.user }}  
db_pass: {{ .Values.dbSettings.password }}  
data_reset: false
```

```
cat gocovid/templates/_helpers.tpl (3)  
...  
{{- define "apiConfig" -}}  
{{ tpl (.Files.Get "conf/config.yaml") . | b64enc }}  
{{- end }}  
...
```

GitLab-ci

Infra

- Pipeline stages
 - init (cache)
 - validate (terraform validate, tflint)
 - build (artefacts: reports, environments)
 - Infracost (only test)
 - apply (test, prod, manual)
 - extra-init (cache)
 - extra-validate (terraform validate, tflint)
 - extra-build (artefacts: reports, environments)
 - extra-apply (test, prod, manual)
 - extra-destroy (PHASE == "destroy")
 - destroy (PHASE == "destroy")
 - Env vars
- Terraform state
- Dependencies
- Needs

Application

- Pipeline stages
 - lint (base)
 - test
 - Sonarqube (only master)
 - build (base, trivy check)
 - helm build (base, environments)
 - deploy (base) (test, prod)
- base build (extends):
 - only branches
 - only protected tags
- Env vars
- Push docker images
- Push helm chart
- Needs to increase

Logs and Monitoring by CloudWatch

Log queries:

```
fields @timestamp, @message, kubernetes.namespace_name,  
kubernetes.pod_name  
| sort @timestamp desc  
| filter kubernetes.pod_name not like /^fluent.*.cz268$/  
| limit 9999
```

```
fields @timestamp, @message, kubernetes.namespace_name,  
kubernetes.pod_name  
| sort @timestamp desc  
| filter kubernetes.namespace_name="gocovid"  
| filter kubernetes.pod_name like /^deploy-gocovid.*/  
| limit 9999
```

```
fields @timestamp, @message, kubernetes.namespace_name,  
kubernetes.pod_name  
| sort @timestamp desc  
| filter kubernetes.pod_name not like /^front.*/  
| limit 9999
```

Monitoring:

- Alarms and SNS
- Build dashboards
- All LoadBalancer Metrics
- All EKS metrics
- All RDS Metrics
- Tags, env tags
- Prometheus
- Grafana

Possible project improvements

- two-factor authentication to gitlab or vpn
- change EFS with S3 bucket on gitlab
- add redis for cache
- generate full sql-query to decrease connection to database
- return errors in main.go (process them in main.go)
- Add .dockerignore file to the context directory
- Describe some tests in Makefile
- Increase tests coverage

Thank you for your attention